



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Algorithm Design and Analysis (H)

CS216

Instructor: Shan CHEN (陈杉)

chens3@sustech.edu.cn

(slides edited from Prof. Shiqi Yu)



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Greedy Algorithms



8. Huffman Codes



Encoding

- **Q.** Why do we need encoding?
- **A.** Encoding transforms data of human language to numbers such that they can be processed by digital computers.
- Ex. Postcode, character codes (Unicode), etc.







- **Q.** Given a text that uses 32 symbols (26 different letters, space, and some punctuation characters), how can we encode this text in bits?
- **A.** We can encode 2^5 different symbols using a fixed length of 5 bits per symbol. This is called **fixed length encoding**.





Fixed Length Encoding

- 64 samples (1 poison)
- How many guinea pigs do we need to find the poison?

						
0	0	0	0	0	0	0
1	0	0	0	0	0	1
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	0	0	1	0	0
...						
63	1	1	1	1	1	1



Data Compression

- **Q.** Some symbols (e, t, a, o, i, n) are used far more often than others. How can we use this to reduce our encoding?
- **A.** Encode such characters with **fewer** bits and the others with more bits.
- **Q.** How do we know when the next symbol begins?
- **A.** Use a separation symbol (like the pause in Morse), or make sure that there is no ambiguity by ensuring that **no** code is a **prefix** of another one.
- **Ex.** $c(a) = 01$, $c(b) = 010$, $c(e) = 1$.
- **Q.** What is 0101?





Prefix Codes

- **Def.** A **prefix code** for a set S is a function γ that maps each $x \in S$ to a bit string such that for $x, y \in S$, $x \neq y$, $\gamma(x)$ is **not a prefix** of $\gamma(y)$.
- **Ex.** $c(a) = 11$, $c(e) = 01$, $c(k) = 001$, $c(l) = 10$, $c(u) = 000$.
- Q. What is the meaning of 1001000001?
- A. "leuk"
- **Ex.** Frequencies in 1G text: $f_a = 0.4$, $f_e = 0.2$, $f_k = 0.2$, $f_l = 0.1$, $f_u = 0.1$.
- Q. What is the size of the encoded text?
- A. $2f_a + 2f_e + 3f_k + 2f_l + 3f_u = 2.3G$





Optimal Prefix Codes

- **Def.** The **average bits per letter (ABL)** of a prefix code γ is the sum, over all symbols $x \in S$, of its frequency f_x times the number of bits of its encoding $\gamma(x)$:

$$ABL(\gamma) = \sum_{x \in S} f_x |\gamma(x)|$$


- **Q.** Can we construct a prefix code that has the lowest ABL?
- **Observation.** A prefix code can be modeled in a binary tree.
 - Any binary code can be modeled in a binary tree.





Representing Prefix Codes with Binary Trees


- **Ex.** $c(a) = 11$, $c(e) = 01$, $c(k) = 001$, $c(l) = 10$, $c(u) = 000$.
- **Q.** How does the tree of a prefix code look like?
- **A.** Only the **leaves** have a label.
- **Pf.** An encoding of x is a prefix of an encoding of y if and only if the path of x is a prefix of the path of y .





Representing Prefix Codes with Binary Trees

- **Q.** What is the meaning of 111010001111101000?
- **A.** "simplel"




- **Q.** How can this prefix code be made more efficient?
- **A.** Change encoding of s and p to a **shorter** one.





Representing Prefix Codes with Binary Trees

- **Def.** A tree is **full** if every node that is not a leaf has two children.
- **Claim.** The binary tree corresponding to an **optimal** prefix code is full.
- **Pf. (by contradiction)**
 - Suppose T is binary tree of optimal prefix code and is not full.
 - This means there is a node u with only one child v .
 - Case 1: u is the root.
 - ✓ Delete u and use v as the root.
 - Case 2: u is not the root.
 - ✓ Delete u and make v be a child of w in place of u . (w is the parent of u .)
 - In both cases the number of bits needed to encode any leaf in the subtree of v is **decreased**. The rest of the tree is not affected.
 - Clearly this new tree T' has a smaller ABL than T . Contradiction!





Optimal Prefix Codes: False Start



- **Q.** Where in the tree of an optimal prefix code should letters with a high frequency be placed?
- **A.** Near the top.
- **Greedy template.** [Shannon-Fano, 1949] Create tree **top-down**. Split S into two sets S_1 and S_2 with (almost) equal frequencies. Recursively build tree for S_1 and S_2 .
- **Q.** Does this approach output an optimal prefix code?

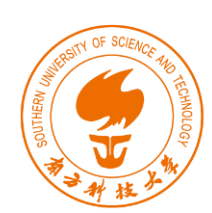




Optimal Prefix Codes: False Start


- **Greedy template.** [Shannon-Fano, 1949] Create tree **top-down**. Split S into two sets S_1 and S_2 with (almost) equal frequencies. Recursively build tree for S_1 and S_2 .
- **Counterexample:** $f_a = 0.32$, $f_e = 0.25$, $f_k = 0.20$, $f_l = 0.18$, $f_u = 0.05$





Optimal Prefix Codes: Huffman Encoding

- **Observation.** Lowest frequency items should be at the lowest level in the tree of an optimal prefix code.
- **Observation.** For $n > 1$, the lowest level always contains at least two leaves.
- **Observation.** The order in which items appear in a level does not matter.
- **Claim.** There is an optimal prefix code with tree T^* where the **two lowest-frequency letters** are assigned to **leaves that are siblings** in T^* .
- **Greedy template.** [Huffman, 1952] Create tree **bottom-up**. Make two leaves for two lowest-frequency letters y and z . Recursively build tree for the rest using a meta-letter for y, z .





Optimal Prefix Codes: Huffman Encoding

```
Huffman(S) {  
  if |S| == 2 {  
    return tree with root and 2 leaves  
  } else {  
    let y and z be lowest-frequency letters in S  
    let S' be S with y and z removed  
    insert new letter  $\omega$  in S' with  $f_{\omega} = f_y + f_z$   
    T' = Huffman(S')  
    T = add two children y and z to leaf  $\omega$  from T'  
    return T  
  }  
}
```

- **Q.** What is the time complexity?
- **A.** $T(n) = T(n - 1) + O(\log n)$, therefore $T(n) = O(n \log n)$.

↖ ExtractMin from priority queue S





Huffman Encoding: Greedy Analysis

- **Claim.** $ABL(T) = ABL(T') + f_\omega$.
- Pf. Recall that $T = T'$ with two children x, y added to leaf ω .

$$\begin{aligned} ABL(T) &= \sum_{x \in S} f_x \cdot \text{depth}_T(x) \\ &= f_{y^*} \cdot \text{depth}_T(y^*) + f_{z^*} \cdot \text{depth}_T(z^*) + \sum_{x \neq y^*, z^*} f_x \cdot \text{depth}_T(x) \\ &= (f_{y^*} + f_{z^*}) \cdot (1 + \text{depth}_{T'}(\omega)) + \sum_{x \neq y^*, z^*} f_x \cdot \text{depth}_{T'}(x) \\ &= f_\omega \cdot (1 + \text{depth}_{T'}(\omega)) + \sum_{x \neq y^*, z^*} f_x \cdot \text{depth}_{T'}(x) \\ &= f_\omega + f_\omega \cdot \text{depth}_{T'}(\omega) + \sum_{x \neq y^*, z^*} f_x \cdot \text{depth}_{T'}(x) \\ &= f_\omega + \sum_{x \in S'} f_x \cdot \text{depth}_{T'}(x) \\ &= f_\omega + ABL(T'). \quad \blacksquare \end{aligned}$$





Huffman Encoding: Greedy Analysis

- **Claim.** Huffman code for S achieves the **minimum ABL** of any **prefix code**.
- Pf. (by induction over $n = |S|$)
 - **Basis Step:** For $n = 2$ there is no shorter code than root and two leaves.
 - **Inductive Hypothesis (IH):** Suppose Huffman tree T' for S' of size $n - 1$, with leaf ω instead of the lowest-frequency nodes y and z , is optimal.
 - **Inductive Step:** (by contradiction) [proof idea]
 - ✓ Suppose some other tree Z of size n is better.
 - ✓ Delete lowest frequency items y and z from Z creating Z' .
 - ✓ Z' cannot be better than T' by IH.





Huffman Encoding: Greedy Analysis

- **Claim.** Huffman code for S achieves the **minimum ABL** of any **prefix code**.
- Pf. (by induction over $n = |S|$)
 - **Basis Step:** For $n = 2$ there is no shorter code than root and two leaves.
 - **Inductive Hypothesis (IH):** Suppose Huffman tree T' for S' of size $n - 1$, with leaf ω instead of the lowest-frequency nodes y and z , is optimal.
 - **Inductive Step: (by contradiction)**
 - ✓ Let T be the Huffman tree. Suppose there is tree Z such that $ABL(Z) < ABL(T)$.
 - ✓ Then w.l.o.g. we assume Z has leaves y and z as siblings (see Observation).
 - ✓ Let Z' be Z with y and z deleted and their former parent labeled ω' .
 - ✓ We know that $ABL(Z) = ABL(Z') + f_{\omega'}$, as well as $ABL(T) = ABL(T') + f_{\omega}$.
 - ✓ Recall $ABL(Z) < ABL(T)$ and $f_{\omega'} = f_{\omega}$, so $ABL(Z') < ABL(T')$. Contradiction with IH!





Huffman Encoding: Summary

- **Greedy approach:** Shrink the size of the problem instance, so that an equivalent smaller problem can then be solved by recursion.
 - The greedy operation is proved to be "safe": solving the smaller problem still leads to an optimal solution for the original problem.
- **Application.** ZIP file format that supports lossless data compression.
 - Its most common compression algorithm DEFLATE uses LZ77 and Huffman.
 - Document: <https://pkware.cachefly.net/webdocs/APPNOTE/APPNOTE-6.2.0.txt>

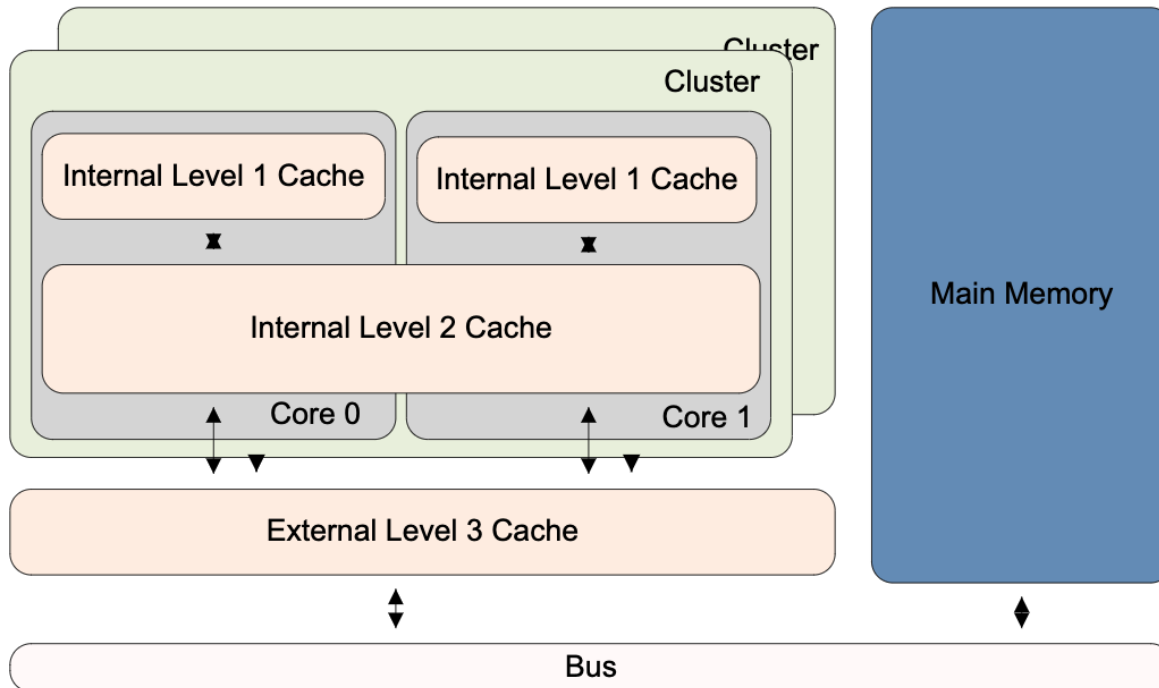


9. Optimal Caching



Caching

- HD->memory->cache



Memory type	Typical size	Typical access time
Processor registers	128KB	1 cycle
On-chip L1 cache	32KB	1-2 cycle(s)
On-chip L2 cache	128KB	8 cycles
Main memory (L3) dynamic RAM	MB or GB ^[1]	30-42 cycles
Back-up memory (hard disk) (L4)	MB or GB	> 500 cycles

^[1] Size limited by the processor core addressing, for example a 32-bit processor without memory management can directly address 4GB of memory.



Optimal Offline Caching

- **Offline caching.**

- Cache with capacity to store k items.
- Sequence of m item **requests** d_1, d_2, \dots, d_m . ← offline: known sequence
- Cache **hit**: item already in cache when requested.
- Cache **miss**: item not already in cache when requested.
 - ✓ Must bring requested item into cache, and evict some existing item if full.

- **Goal.** Eviction schedule that **minimizes** number of **evictions**.

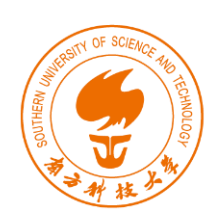
- **Ex:** $k = 2$, initial cache = ab , requests: a, b, c, b, c, a, a, b .
- Optimal eviction schedule: 2 cache misses.



Optimal Offline Caching

- **Greedy templates.**

- LIFO. (last-in-first-out) Evict item brought in least recently.
- FIFO. (first-in-first-out) Evict item brought in most recently.
- LRU. (least-recently-used) Evict item whose most recent access was earliest.
- LFU. (least-frequently-used) Evict item that was least frequently requested.



Optimal Offline Caching

- Greedy templates.

- LIFO. (last-in-first-out)
- FIFO. (first-in-first-out)
- LRU. (least-recently-used)
- LFU. (least-frequently-used)

		cache							
requests	⋮			
	<i>a</i>	<i>a</i>	<i>w</i>	<i>x</i>	<i>y</i>	<i>z</i>		FIFO: eject <i>a</i>	
	<i>d</i>	<i>a</i>	<i>w</i>	<i>x</i>	<i>d</i>	<i>z</i>		LRU: eject <i>d</i>	
	<i>a</i>	<i>a</i>	<i>w</i>	<i>x</i>	<i>d</i>	<i>z</i>			
	<i>b</i>	<i>a</i>	<i>b</i>	<i>x</i>	<i>d</i>	<i>z</i>			
	<i>c</i>	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>z</i>			
	<i>e</i>	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>		LIFO: eject <i>e</i>	
	<i>g</i>	?	?	?	?	?			
	<i>b</i>								
	<i>e</i>								
	<i>d</i>								
	⋮								

cache miss
(which item to eject?)



- current cache:


a	b	c	d	e	f
---	---	---	---	---	---

- **Theorem.** [Bellady, 1960s] FF is optimal eviction schedule.
- Algorithm and theorem are intuitive; proof is subtle (shown later).



Reduced Eviction Schedules

- **Def.** A **reduced** schedule is a schedule that only inserts an item into the cache in a step when that item is **requested** and **not yet in cache**.





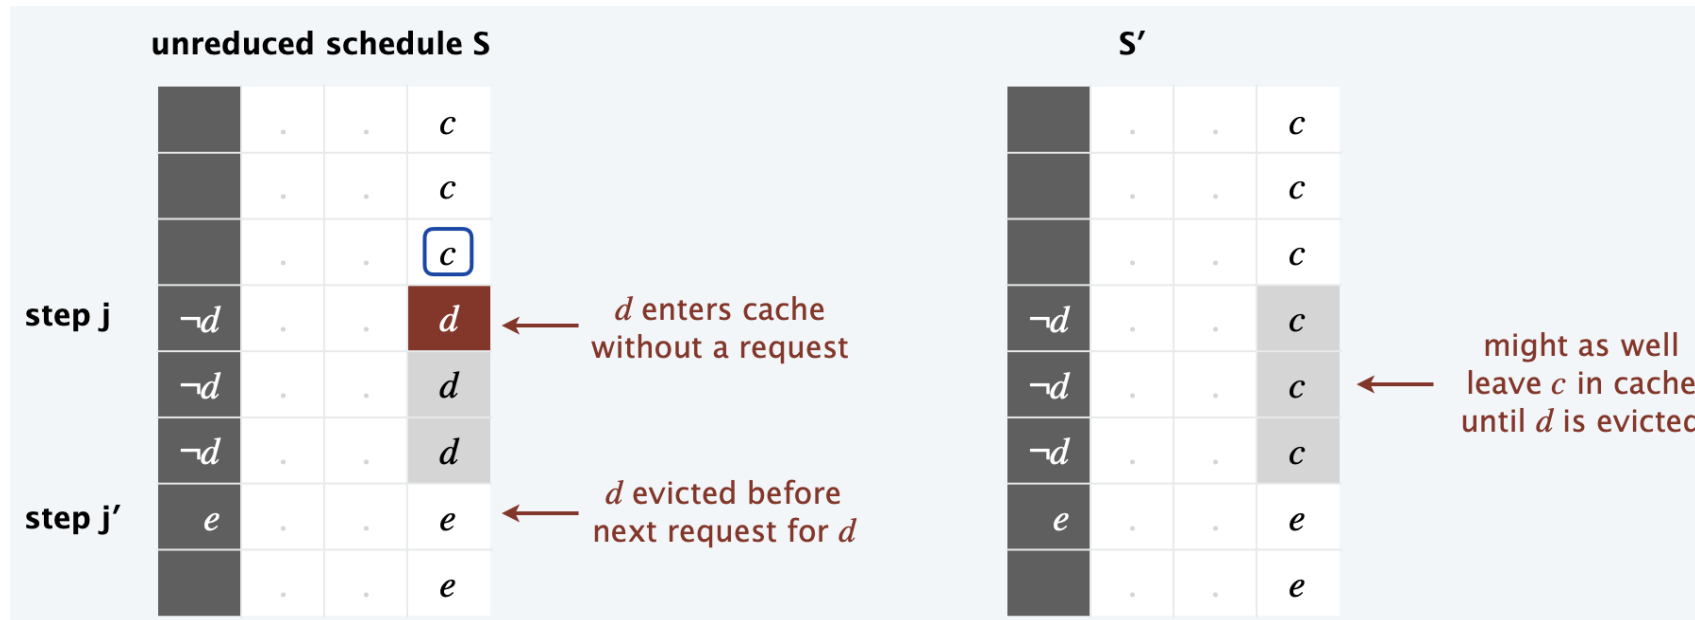
Reduced Eviction Schedules

- **Claim.** Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.
- Pf. (by induction on number of steps j) [proof idea]
 - Base case: $j = 0$.
 - Case 1: S brings d into the cache in step j without a request.
 - Case 2: S brings d into the cache in step j even though d is in cache.
 - If multiple unreduced items in step j , apply each one in turn.
 - Deal with Case 1 before Case 2, as resolving Case 1 might trigger Case 2 ■



Reduced Eviction Schedules

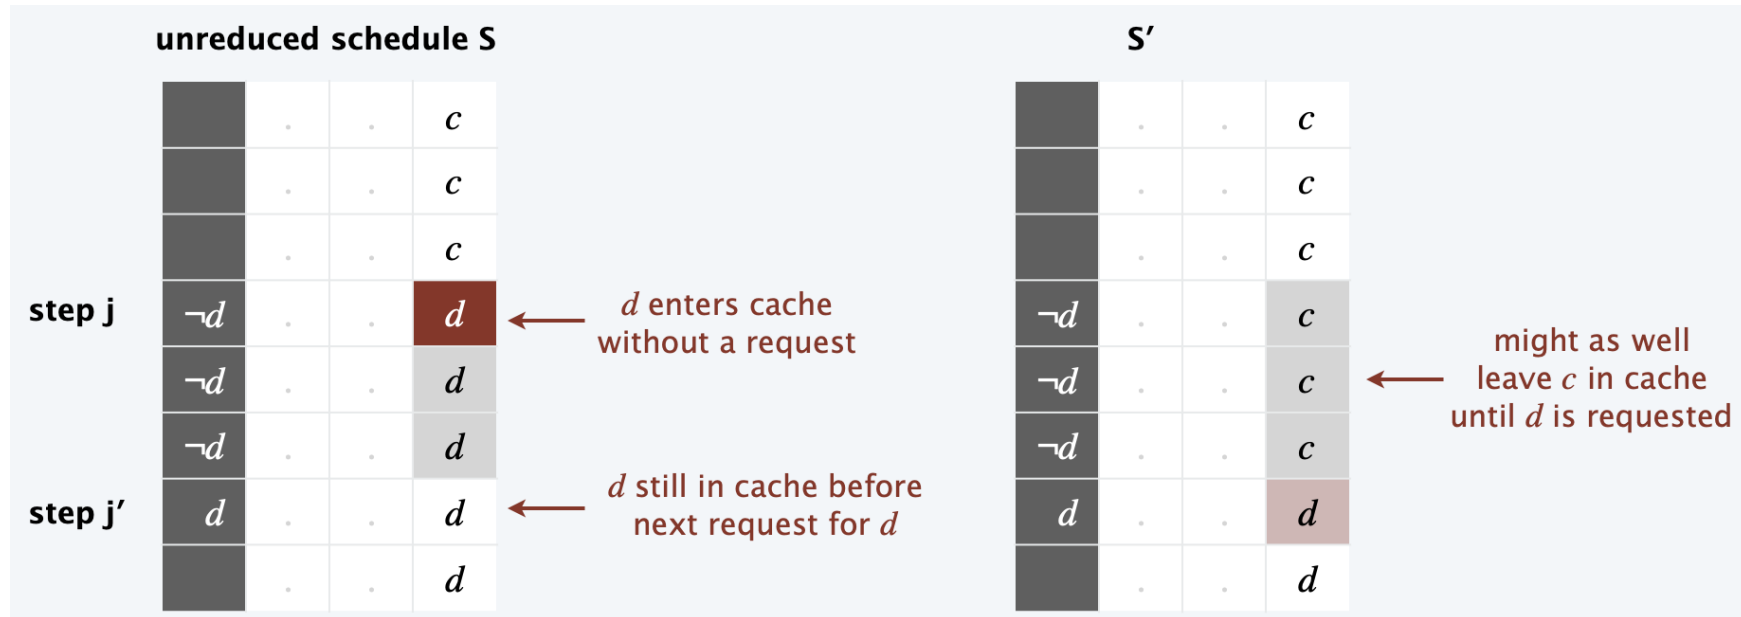
- **Claim.** Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.
- Pf. (by induction on number of steps j) [d is inserted in step j]
 - Let c be the item S evicts when it brings d into the cache.
 - Case 1a: d evicted before next request for d .

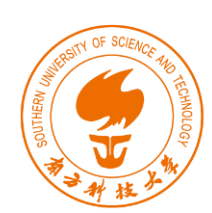




Reduced Eviction Schedules

- **Claim.** Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.
- Pf. (by induction on number of steps j) [d is inserted in step j]
 - Let c be the item S evicts when it brings d into the cache.
 - Case 1b: next request for d occurs before d is evicted.





Reduced Eviction Schedules

- **Claim.** Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.
- Pf. (by induction on number of steps j) [d is inserted in step j]
 - Let c be the item S evicts when it brings d into the cache.
 - Case 2a: d evicted before it is needed.

unreduced schedule S					S'				
step j		d_1	a	c			d_1	a	c
		d_1	a	c			d_1	a	c
		d_1	a	c			d_1	a	c
	d	d_1	a	d_3	←	d	d_1	a	c
	d	d_1	a	d_3	←	d	d_1	a	c
step j'	c	c	a	d_3		c	c	a	c
	b	c	a	b	←	b	c	a	b
	d	c	a	d_3	←	d	c	a	d_3

d_3 enters cache even though d_1 is already in cache

d_3 not needed

d_3 evicted

d_3 needed

might as well leave c in cache until d_3 is evicted



Reduced Eviction Schedules

- **Claim.** Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.
- Pf. (by induction on number of steps j) [d is inserted in step j]
 - Let c be the item S evicts when it brings d into the cache.
 - Case 2b: d needed before it is evicted.

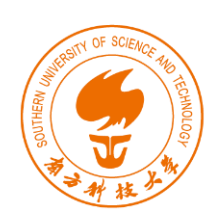
unreduced schedule S					S'				
step j		d_1	a	c			d_1	a	c
		d_1	a	c			d_1	a	c
		d_1	a	c			d_1	a	c
	d	d_1	a	d_3	←	d	d_1	a	c
	d	d_1	a	d_3	←	d	d_1	a	c
	c	c	a	d_3		c	c	a	c
step j'	a	c	a	d_3		a	c	a	c
	d	c	a	d_3	←	d	c	a	d_3

Annotations for unreduced schedule S:

- At step j , c is boxed in the cache.
- At step j , d_3 enters cache even though d_1 is already in cache.
- At step j , d_3 not needed.
- At step j' , d_3 needed.

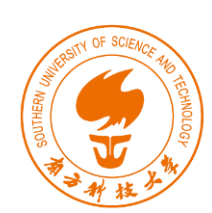
Annotation for S':

- At step j , c is not in the cache; d_3 is in the cache. Annotation: might as well leave c in cache until d_3 is needed.



Farthest-In-Future: Greedy Analysis

- **Theorem.** Farthest-in-future (FF) is optimal eviction algorithm.
- Pf. Follows directly from the following invariant.
- **Invariant.** There exists an optimal reduced schedule S that has the same eviction schedule as S_{FF} through the first j steps.
- Pf. (by induction on number of steps j)
 - **Basis Step:** $j = 0$.
 - **Inductive Step:** Let S be reduced schedule that satisfies invariant through j steps. We produce S' that satisfies invariant after $j + 1$ steps.
 - ✓ Let d denote the item requested in step $j + 1$.
 - ✓ Since S and S_{FF} have agreed up until now, they have the same cache contents before step $j + 1$.



Farthest-In-Future: Greedy Analysis

- **Theorem.** Farthest-in-future (FF) is optimal eviction algorithm.
- Pf. Follows directly from the following invariant.
- **Invariant.** There exists an optimal reduced schedule S that has the same eviction schedule as S_{FF} through the first j steps.
- Pf. (by induction on number of steps j) [d is requested in step $j + 1$]
 - **Inductive Step:** Let S be reduced schedule that satisfies invariant through j steps. We produce S' that satisfies invariant after $j + 1$ steps.
 - Case 1: d is already in the cache.
 - ✓ $S' = S$ satisfies invariant.
 - Case 2: d is not in the cache and S and S_{FF} evicts the same item.
 - ✓ $S' = S$ satisfies invariant.



Farthest-In-Future: Greedy Analysis

- Pf. (continued)

- Case 3: d is not in the cache; S_{FF} evicts e ; S evicts $f \neq e$.
 - ✓ Begin construction of S' from S by evicting e instead of f .

j	same	e	f	same	e	f
	S			S'		
j+1	same	e	d	same	d	f
	S			S'		

- ✓ Now S' agrees with S_{FF} on first $j + 1$ requests; we show that having f in cache is no worse than having e in cache. must involve e or f or both ↘
- ✓ Let S' behave the same as S until S' is forced to take a different action, because (Case 3a/3b) either e or f is requested or because (Case 3c) S evicts e .



Farthest-In-Future: Greedy Analysis

- Pf. (continued) [Let j' be the first step after $j + 1$ that S and S' must take different actions, and let g be the item requested at step j' .]

- Case 3: d is not in the cache; S_{FF} evicts e ; S evicts $f \neq e$; S' evicts e .



S

S'

S' agrees with S_{FF} through step $j+1$

- Case 3a: $g = e$.
 - ✓ Can't happen with FF since there must be a request for f before e .
- Case 3b: $g = f$.
 - ✓ Element f is not in cache of S , so let e' be the element that S evicts.
 - ✓ If $e' = e$, S' accesses f from cache; now S and S' have same cache.
 - ✓ If $e' \neq e$, S' evicts e' and brings e into cache; now S and S' have the same cache.
 - ✓ Let S' behave exactly like S for remaining requests.

S' is no longer reduced, but can be transformed into a reduced schedule that agrees with S_{FF} through step $j+1$



Farthest-In-Future: Greedy Analysis

- Pf. (continued) [Let j' be the **first** step after $j + 1$ that S and S' must take different actions, and let g be the item requested at step j' .]
 - Case 3: d is not in the cache; S_{FF} evicts e ; S evicts $f \neq e$; S' evicts e .

j'	same	e	same	f
------	------	-----	------	-----

S

S'

- Case 3c: S evicts e (and $g \neq e, f$).
 - ✓ Make S' evict f .

j'	same	g	same	g
------	------	-----	------	-----

S

S'

- ✓ Now S and S' have the same cache.
- ✓ Let S' behave exactly like S for the remaining requests. ▀



Optimal Caching: Summary

- **Online vs. offline caching.**

- Offline: full sequence of requests is known a priori.
- Online (reality): requests are not known in advance.
- Caching is among most fundamental online problems in CS.

- **LRU.** Evict item whose most recent access was earliest.

FF with direction of time reversed!

- **Theorem.** FF is optimal offline eviction algorithm.

- Provides basis for understanding and analyzing online algorithms.
- LIFO can be arbitrarily bad.
- LRU is k -competitive, i.e., for any sequence of requests R , $\text{LRU}(R) \leq k \text{FF}(R) + k$.
- Randomized marking is $O(\log k)$ -competitive.

shown later in Section 13.8