



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Algorithm Design and Analysis (H)

CS216

Instructor: Shan CHEN (陈杉)

chens3@sustech.edu.cn

(slides edited from Prof. Shiqi Yu)



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Dynamic Programming




6. Shortest Paths with Negative Weights



Shortest Paths with Negative Weights


- **Shortest-path problem.** Given a digraph $G = (V, E)$, with **arbitrary** edge weights c_{vw} , find shortest path from source node s to destination node t .
assume there exists a path from every node to t
- **Ex.** Nodes represent agents in a financial setting and c_{vw} is cost of transaction in which we buy from agent v and sell immediately to w .






Shortest Paths: Failed Attempts

- **Dijkstra.** Can fail if there exist negative edge weights.




- **Re-weighting.** Adding a constant to every edge weight can still fail.







Shortest Paths: Negative Cycles

- **Negative cycles.** A **negative cycle** is a directed cycle for which the sum of its edge lengths is negative.



- **Observation.** If some s - t path contains a negative cycle, then there does not exist a shortest s - t path; if there exists no negative cycle, there exists a shortest s - t path that is simple (and has $\leq n - 1$ edges).





Shortest-Paths and Negative-Cycle Problems

- **Single-destination shortest-paths problem.** Given a digraph $G = (V, E)$, with **arbitrary** edge weights c_{vw} (but **no negative cycles**) and a destination node t , find v - t shortest paths **from every node v** to t .

Single-destination shortest-paths problem is equivalent to **single-source** shortest-paths problem with edge directions reversed.

- **Negative cycle detection problem.** Given a digraph $G = (V, E)$, with **arbitrary** edge weights c_{vw} , find a negative cycle (if one exists).



Shortest Paths: Dynamic Programming

- **Def.** $\text{OPT}(i, v)$ = length of shortest v - t path P using $\leq i$ edges.
- **Goal.** $\text{OPT}(n - 1, v)$ ← if no neg cycles, there exists a simple shortest path
- **To compute $\text{OPT}(i, j)$:**
 - **Case 1:** P uses at most $\leq i - 1$ edges.
 - ✓ $\text{OPT}(i, v) = \text{OPT}(i - 1, v)$
 - **Case 2:** P uses exactly i edges.
 - ✓ let (v, w) be the first edge in P : pay the cost of c_{vw} , then select best w - t path using $\leq i - 1$ edges



Shortest Paths: Dynamic Programming

- **Def.** $OPT(i, v)$ = length of shortest v - t path P using at most i edges.
- **Goal.** $OPT(n - 1, v)$ ← if no neg cycles, there exists a simple shortest path
- **Bellman equation.**

$$OPT(i, j) = \begin{cases} 0, & \text{if } i = 0 \text{ and } v = t \\ \infty, & \text{if } i = 0 \text{ and } v \neq t \\ \min\{OPT(i - 1, v), \min_{w \in V} (OPT(i - 1, w) + c_{vw})\} & \text{if } i > 0 \end{cases}$$



Shortest Paths: Algorithm


- **Dynamic programming algorithm (bottom-up).**

```
Shortest-Path(G, t) {  
    foreach node  $v \in V$   
         $M[0, v] = \infty$   
     $M[0, t] = 0$   
  
    for  $i = 1$  to  $n - 1$   
        foreach node  $v \in V$   
             $M[i, v] = M[i - 1, v]$   
            foreach edge  $(v, w) \in E$   
                 $M[i, v] = \min\{ M[i, v], M[i - 1, w] + c_{vw} \}$   
}
```

- **Finding shortest paths.** Maintain *successor*[i, v] for each $M[i, v]$.



Shortest-Paths Algorithm: Demo



(a)

	0	1	2	3	4	5
t	0	0	0	0	0	0
a	∞	-3	-3	-4	-6	-6
b	∞	∞	0	-2	-2	-2
c	∞	3	3	3	3	3
d	∞	4	3	3	2	0
e	∞	2	0	0	0	0

(b)

$$\text{OPT}(i, j) = \begin{cases} 0, & \text{if } i = 0 \text{ and } v = t \\ \infty, & \text{if } i = 0 \text{ and } v \neq t \\ \min\{\text{OPT}(i-1, v), \min_{w \in V} (\text{OPT}(i-1, w) + c_{vw})\} & \text{if } i > 0 \end{cases}$$



Shortest Paths: Algorithm

- **Dynamic programming algorithm (bottom-up).**

```
Shortest-Path(G, t) {  
  foreach node  $v \in V$   
     $M[0, v] = \infty$   
   $M[0, t] = 0$   
  
  for  $i = 1$  to  $n - 1$   
    foreach node  $v \in V$   
       $M[i, v] = M[i - 1, v]$   
      foreach edge  $(v, w) \in E$   
         $M[i, v] = \min\{ M[i, v], M[i - 1, w] + c_{vw} \}$   
}
```

only $M[i - 1, .]$ is used!

- **Running time.** $O(mn)$ **Space.** $O(n^2)$ space

can we reduce the required space?



Shortest Paths: Practical Improvements

- **Space optimization.** Maintain two 1-D arrays (instead of 2-D array).
 - $d[v]$ = length of a shortest v - t path that we have found so far
 - $successor[v]$ = next node on a v - t path
- **Performance optimization.** If $d[w]$ was not updated in iteration $i - 1$, then no need to consider edges entering w in iteration i .



Bellman-Ford-Moore: Efficient Implementation

- **Dynamic programming algorithm (bottom-up).**

```
Bellman-Ford-Moore(G, s, t) {  
    foreach node v ∈ V  
        d[v] = ∞  
        successor[v] = null  
  
    d[t] = 0  
    for i = 1 to n - 1  
        foreach node w ∈ V  
            if (d[w] has been updated in previous iteration) {  
                foreach node v such that (v, w) ∈ E {  
                    if (d[v] > d[w] + cvw)  
                        d[v] = d[w] + cvw  
                        successor[v] = w  
                }  
            }  
        if (no d[w] value changed in iteration i)  
            break  
}
```

O(n) space

push-based rather than pull-based

*each pass
O(m) time:
O(mn) total*



Bellman-Ford-Moore: Analysis

- **Theorem.** After pass i , $d[v]$ = length of a shortest v - t path using $\leq i$ edges.
- **Pf. (by induction on i)**
 - **Base case:** $i = 0$.
 - **Inductive case:** Assume true after pass i . Let P be any v - t path with $\leq i + 1$ edges.
 - ✓ Let (v, w) be first edge in P and let P' be subpath from w to t .
 - ✓ By inductive hypothesis, because P' is a w - t path with $\leq i$ edges, at the end of pass i we have $d[w] \leq \ell(P')$.
 - ✓ After considering edge (v, w) in pass $i + 1$ (or in some previous pass $< i + 1$):
 $d[v] \leq c_{vw} + d[w]$. Then, since $d[w] \leq \ell(P')$ after pass i and $d[w]$ never increases during the algorithm, we have $d[v] \leq c_{vw} + \ell(P') = \ell(P)$.
 - ✓ Obviously, there exists a v - t path of length $d[v]$. From above, this path is a shortest v - t path using $\leq i + 1$ edges.



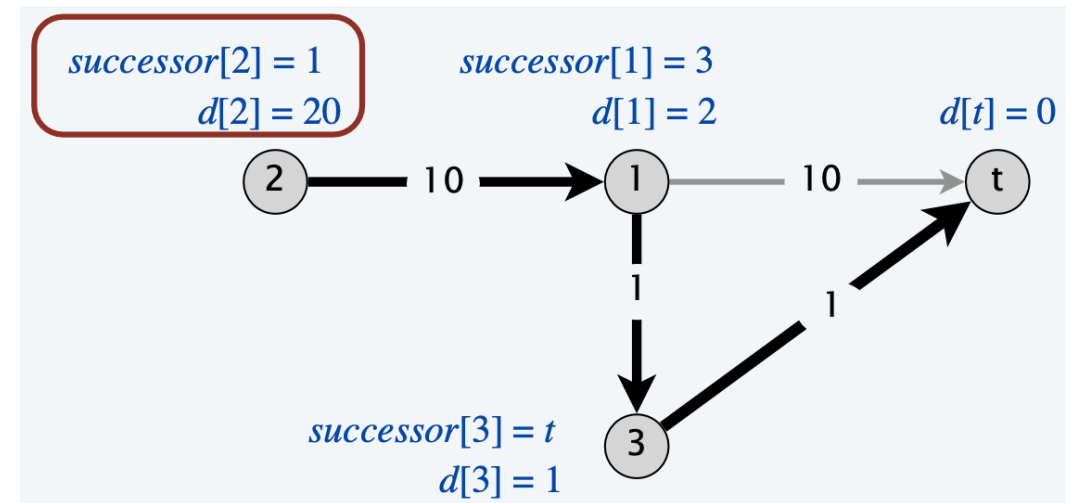
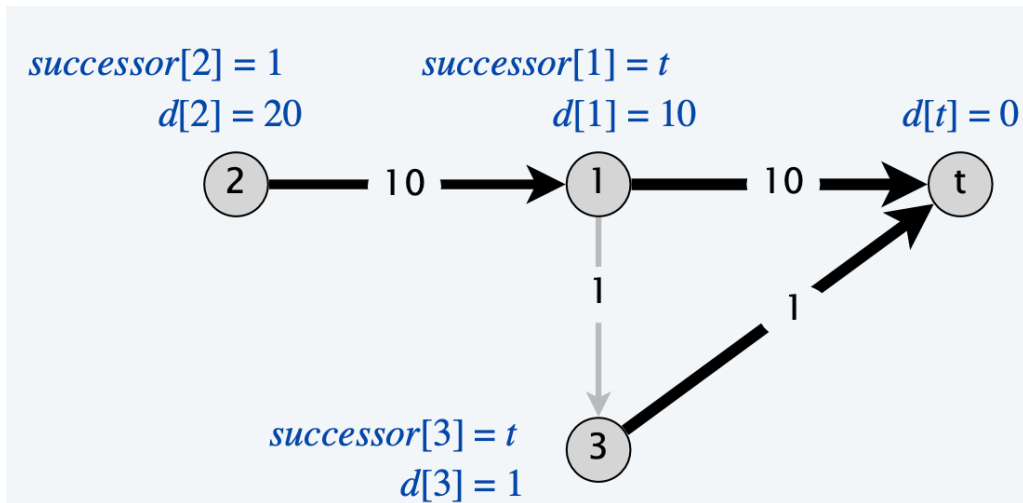
Bellman-Ford-Moore: Analysis

- **Theorem.** Assume **no negative cycles**, Bellman-Ford-Moore computes the lengths of the shortest v - t paths in $O(mn)$ time and $O(n)$ extra space.
- **Pf.** From previous observation and theorem, we have:
 - If no negative cycles, shortest path exists and has at most $n - 1$ edges.
 - After pass $n - 1$, $d[v]$ = length of a shortest v - t path using $n - 1$ edges.
- **Remark.** Bellman–Ford–Moore is typically faster in practice.
 - Edge (v, w) is considered in pass $i + 1$ only if $d[w]$ was updated in pass i .
 - If shortest path is known to have k edges, then algorithm finds it in k passes.
- **Q.** How do we find a shortest v - t path of length $d[v]$ for every node v ?



Bellman-Ford-Moore: Finding Shortest Paths



- **Claim.** Throughout Bellman–Ford–Moore, following the $successor[v]$ pointers gives a directed path from v to t of length $d[v]$.
- **Counterexamples.** (Claim is false!)
 - Length of successor v - t path may be **strictly shorter** than $d[v]$.
 - ✓ Ex. Consider nodes in order: $t, 1, 2, 3$





Bellman-Ford-Moore: Finding Shortest Paths

- **Claim.** Throughout Bellman–Ford–Moore, following the *successor*[*v*] pointers gives a directed path from *v* to *t* of length *d*[*v*].
- **Counterexamples.** (Claim is false!)
 - Length of successor *v*-*t* path may be strictly shorter than *d*[*v*].
 - If negative cycles exist, successor graph may have **directed cycles**.
 - ✓ Ex. Consider nodes in order: *t*, 1, 2, 3, 4





Bellman-Ford-Moore: Finding Shortest Paths

- **Lemma.** Any directed cycle W in the **successor graph** is a **negative cycle**.
- **Pf.**
 - If $\text{successor}[v] = w$, we have $d[v] \geq d[w] + c_{vw}$. (They are equal when $\text{successor}[v]$ is set; $d[w]$ can only decrease; $d[v]$ decreases only when $\text{successor}[v]$ is reset.)
 - Let $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k \rightarrow v_1$ be the sequence of nodes in a directed cycle W .
 - Assume that (v_k, v_1) is the last edge in W added to the successor graph.
 - Just prior to that:

$d[v_1]$	\geq	$d[v_2]$	$+$	$\ell(v_1, v_2)$	
$d[v_2]$	\geq	$d[v_3]$	$+$	$\ell(v_2, v_3)$	
\vdots	\vdots	\vdots			
$d[v_{k-1}]$	\geq	$d[v_k]$	$+$	$\ell(v_{k-1}, v_k)$	
$d[v_k]$	$>$	$d[v_1]$	$+$	$\ell(v_k, v_1)$	← holds with strict inequality since we are updating $d[v_k]$
 - Adding inequalities yields $\ell(v_1, v_2) + \ell(v_2, v_3) + \dots + \ell(v_{k-1}, v_k) + \ell(v_k, v_1) < 0$. ▀



Bellman-Ford-Moore: Finding Shortest Paths

- **Theorem.** Assuming no negative cycles, Bellman–Ford–Moore finds shortest v - t paths for every node v in $O(mn)$ time and $O(n)$ extra space.
- **Pf.**
 - From previous lemma, the successor graph cannot have a directed cycle. Thus, following the successor pointers from v yields a directed path to t .
 - Let $v = v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k = t$ be the nodes along this path P .
 - Upon termination, if $\text{successor}[v] = w$, we have $d[v] = d[w] + c_{vw}$. (They are equal when $\text{successor}[v]$ is set; $d[\cdot]$ did not change since algorithm terminates.)

➤ Thus,

$$\begin{array}{rcl} d[v_1] & = & d[v_2] + \ell(v_1, v_2) \\ d[v_2] & = & d[v_3] + \ell(v_2, v_3) \\ \vdots & & \vdots \\ d[v_{k-1}] & = & d[v_k] + \ell(v_{k-1}, v_k) \end{array}$$

➤ Adding equations yields $d[v] = d[t] + \ell(v_1, v_2) + \ell(v_2, v_3) + \dots + \ell(v_{k-1}, v_k)$. ▀



Shortest Paths: Asymptotic Complexity

year	worst case	discovered by
1955	$O(n^4)$	Shimbel
1956	$O(m n^2 W)$	Ford
1958	$O(m n)$	Bellman, Moore
1983	$O(n^{3/4} m \log W)$	Gabow
1989	$O(m n^{1/2} \log(nW))$	Gabow–Tarjan
1993	$O(m n^{1/2} \log W)$	Goldberg
2005	$O(n^{2.38} W)$	Sankowski, Yuster–Zwick
2016	$\tilde{O}(n^{10/7} \log W)$	Cohen–Mądry–Sankowski–Vladu
20xx	???	

single-source shortest paths with weights between $-W$ and W



7. Distance-Vector Protocols



Distance-Vector Routing Protocols

- **Communication network.**

- Node \approx router.
- Edge \approx direct communication link.
- Cost of edge \approx latency of link.

← non-negative costs, but Bellman-Ford-Moore used anyway!

- **Dijkstra's algorithm.** Requires **global information** of network.
- **Bellman-Ford-Moore.** Uses only **local knowledge** of neighboring nodes.
- **Synchronization.** We don't expect routers to run in lockstep. The order in which each edges are processed in Bellman-Ford-Moore is not important. Moreover, algorithm converges even if updates are **asynchronous**.



Asynchronous Shortest-Paths Algorithm

Asynchronous-Shortest-Path(G, s, t)

n = number of nodes in G

Array $M[V]$

Initialize $M[t] = 0$ and $M[v] = \infty$ for all other $v \in V$

Declare t to be active and all other nodes inactive

While there exists an active node

 Choose an active node w

 For all edges (v, w) in any order

$M[v] = \min(M[v], c_{vw} + M[w])$

 If this changes the value of $M[v]$, then

$first[v] = w$

v becomes active

 Endfor

w becomes inactive


 EndWhile

no for loop for nodes in asynchronous version



Distance-Vector Routing Protocols

- **Distance-vector routing protocols.** “routing by rumor”
 - Each router maintains a **vector** of shortest path lengths to **every other node** (distances) and the first hop on each path (directions).
 - Algorithm: each router performs n separate computations, one for each potential destination node.
- **Example applications.** RIP, Xerox XNS RIP, Novell's IPX RIP, Cisco's IGRP, DEC's DNA Phase IV, AppleTalk's RTMP.
- **Caveat.** Edge costs may **change** during algorithm (or fail completely).





Path-Vector Routing Protocols

- **Link-state routing protocols.**

- Each router also stores the **entire path**.
- Based on Dijkstra's algorithm.
- Avoids "counting-to-infinity" problem and related difficulties.
- Requires **significantly more storage**.

not just the distance and first hop

- **Ex.** Border Gateway Protocol (BGP), Open Shortest Path First (OSPF).





8. Negative Cycles



Detecting Negative Cycles


- **Negative cycle detection problem.** Given a directed graph $G = (V, E)$, with **arbitrary** edge weights c_{vw} , find a negative cycle (if one exists).
- **Application. [Currency conversion]** Given n currencies and exchange rates between them, is there an arbitrage opportunity?
 - Fastest algorithm very valuable!





Detecting Negative Cycles


- **Lemma.** If $\text{OPT}(n, v) = \text{OPT}(n - 1, v)$ for every v , then no negative cycles.
- **Pf.** The $\text{OPT}(n, v)$ values have converged \Rightarrow shortest v - t path exists. ▀
- **Lemma.** If $\text{OPT}(n, v) < \text{OPT}(n - 1, v)$ for some node v , then (any) shortest v - t path of length $\leq n$ contains a cycle W . Moreover, W is a negative cycle.
- **Pf.**
 - $\text{OPT}(n, v) < \text{OPT}(n - 1, v) \Rightarrow$ shortest v - t path P has exactly n edges.
 - By pigeonhole principle, the path P must contain a repeated node x .
 - Let W be any cycle in P .
 - Deleting W yields a v - t path with $< n$ edges.
 - Therefore, W is a negative cycle. ▀





Detecting Negative Cycles

- **Theorem.** Can find a negative cycle in $O(mn)$ time and $O(n^2)$ space.
- **Pf.** Add new sink node t and connect all nodes to t with 0-length edges. G has a negative cycle if and only if G' has a negative cycle.
 - **Case 1:** $OPT(n, v) = OPT(n - 1, v)$ for every node v
 - ✓ By previous lemma, there exist no negative cycles.
 - **Case 2:** $OPT(n, v) < OPT(n - 1, v)$ for some node v
 - ✓ Can extract negative cycle from v - t path (cycle cannot contain t since no edge leaves t). ▀





Detecting Negative Cycles

- **Theorem.** Can find a negative cycle in $O(mn)$ time and $O(n)$ extra space.
- **Pf.**
 - Run Bellman–Ford–Moore on G' for $n' = n + 1$ passes (instead of $n' - 1$).
 - If no $d[v]$ values updated in pass n' , then no negative cycles.
 - Otherwise, suppose $d[s]$ updated in pass n' .
 - Define $pass(v)$ = last pass in which $d[v]$ was updated.
 - Observe $pass(s) = n'$, and $pass(v) - 1 \leq pass(successor[v])$ for each v .
 - Following successor pointers ($\geq n'$ edges), we must eventually repeat a node.
 - Previous lemma shows that the corresponding cycle is a negative cycle. ▀
- **Remark.** See textbook for improved version and early termination rule. (Tarjan's subtree disassembly trick.)